

UNIVERSITY OF WATERLOO
Department of Systems Design Engineering

SYDE 462: Spring Term Final Report

Group 2: Relay
Adaptive Traffic Control Framework

prepared by

Alex Huras – 20344660
D. Scott Neil – 20349210
Myles Tan – 20349217
Riley Donelson – 20342815

Table of Contents

1	Background Information	1
2	Engineering Design	2
1.0	Front End	2
1.1	Design Research	2
1.2	Design Process	2
1.2.1	Wireframing	2
1.2.2	Medium-Fidelity Prototyping	5
1.2.3	High-Fidelity Prototyping	9
1.3	Application Architecture	17
1.3.1	Framework Selection	18
1.3.2	Mapping Library Selection	21
1.3.3	Visualization Library Selection	23
1.3.4	Architecture Development	26
2.0	Back End	29
3.0	Deep End	29
3.1	Learning in Relay	29
3.1.1	Design Parameter Updates	29
3.1.2	Time Delay Estimation	30
3.1.3	Updating Intra-Agent Behaviour Graphs	30
3.1.3.1	Behaviour Probability Matrix (BPM)	30
3.1.3.2	Learning Probabilities	31
3.1.3.3	Updating BPMs	31
3	Design Testing and Validation	33

1.0	Front End	33
2.0	Back End	33
3.0	Deep End	33
4	Recommended Design Modifications	34
5	Timeline and Project Management	35
1.0	Conclusion	36

List of Figures

Figure 2.1 — Initial wireframe for histogram approach.	3
Figure 2.2 — Application wireframes showing interaction controls and pop-up menu.	3
Figure 2.3 — Sketch of potential side menu for the Relay Interface application.	4
Figure 2.4 — Sketch of a detailed pop-up menu showing intersection data.	4
Figure 2.5 — Balsamiq mockup of a data layer in the Relay app.	5
Figure 2.6 — Balsamiq mockup of the Relay info box.	6
Figure 2.7 — Balsamiq mockup of the Relay dashboard.	7
Figure 2.8 — Balsamiq mockup of the Intersections page.	8
Figure 2.9 — Balsamiq mockup of the Network page.	9
Figure 2.10 — High-fidelity rendering of the Relay branding.	10
Figure 2.11 — 16 pixel grid foundation for the Relay user interface.	11
Figure 2.12 — An intersection name set in 24pt Helvetica Neue Bold.	11
Figure 2.13 — A chart showing the flow of traffic through an intersection over time.	12
Figure 2.14 — A screenshot of the intersection Status data layer.	13
Figure 2.15 — A screenshot of the flow per-intersection heat map.	13
Figure 2.16 — A screenshot of the flow between intersections heat map.	14
Figure 2.17 — A screenshot of the average queue length at each intersection.	15
Figure 2.18 — A screenshot of the connection status window.	16
Figure 2.19 — A screenshot of the Relay intersection dashboard.	17
Figure 2.20 — An architectural diagram for Iteration 1.	27
Figure 2.21 — An architectural diagram for Iteration 2.	28
Figure 2.22 — An architectural diagram for Iteration 3.	29
Figure 2.23 — Example time delay estimation.	31
Figure 2.24 — Example intra-agent edge probabilities for node 0.	32

Glossary of Terms

Angular.js A client-side JavaScript Framework.

Backbone.js A client-side JavaScript Framework.

Chart.js A simple, static, front-end visualization library.

D3.js D3, Data Driven Documents, is a rich, static, front-end data visualization framework.

DOJO A client-side Javascript Framework.

Flot Flot is a JavaScript-based front-end data visualization framework which emphasizes 2D plots and time-series data.

JavaScript MVC A client-side JavaScript Framework.

jQuery A client-side JavaScript Utility Library.

MVC A software development pattern which categorizes classes into three categories: model, view, and controller. Generally, models are responsible for data and business logic, views are responsible for presenting information to users and collecting it, and controllers are responsible for handling interaction between the model and the view.

MV Another software development pattern. MV is similar to MVC, however there is no controller, as controller functionality is built into the view class. This pattern is common in less complex solutions or in languages which do not support a rigid class-based system.

Rickshaw A static front-end visualization library build on the D3 framework.

Underscore.js A client-side JavaScript Utility Library.

Chapter 1

Background Information

Chapter 2

Engineering Design

1.0 Front End

1.1 Design Research

1.2 Design Process

With sufficient Design Research conducted, and functional requirements compiled, the front-end application design moved into an iterative process of defining the user interface (UI), interactions, and visual design of the app through various levels of fidelity. The use of many design tools was employed, as best suited for each given stage of the design process. These tools include Adobe Photoshop and Illustrator, HTML/CSS/JavaScript, and sketching using pencil and paper to quickly illustrate low-fidelity ideas.

1.2.1 Wireframing

The process began with a wireframing stage, involving the quick iteration of information architecture and UI techniques, to discover optimal implementations for each of the functional requirements. Global app navigation, page layout, modularity, and fundamental UI elements were initiated at this stage, to begin to define the design language used in the application. This was accomplished mainly through static paper-based sketches of various screens. This allowed for rapid generation of concepts, and movement between and

through ideas.




Figure 2.1: Initial wireframe for histogram approach.




Figure 2.2: Application wireframes showing interaction controls and pop-up menu.

As can be seen in Figures 2.1 - 2.4, various overlay techniques using both histograms and circular polygons on a map of a city were explored here in this early stage. The histogram approach sought to model each vehicle (or small group of vehicles) on the road as a probability density function that could be moved along a road based on knowledge of traffic presence at each intersection, and speed limit data on each road. More vehicles means more histograms, shown as translucent graphs which when overlaid, become more and more opaque, thus indicating a higher traffic density in that area.

Response for this concept was generally critical, as probability functions overlaid on a map were found to not be particularly easy to understand, at least not immediately at a glance. It was evident at this stage that a simpler - more intuitive - visualization was needed for this application to be useful for both consumers and traffic engineers.

Through exploration of the circular polygon heat-map concept, a more agreeable solution was found. This concept modelled traffic data on a per intersection basis. In other words, when an intersection is experiencing a high volume of traffic, a translucent circular polygon is triggered and overlaid onto a map of the city at the longitude and latitude of the intersection. The size of the circle is positively correlated to the




Figure 2.3: Sketch of potential side menu for the Relay Interface application.




Figure 2.4: Sketch of a detailed pop-up menu showing intersection data.

volume of traffic at the intersection. A second degree of data is then shown through use of colour, to illustrate the performance of the intersection itself. This gives insight into how well the intersection is moving traffic given it's high-volume state. It was agreed upon that a well-performing high-volume intersection should receive a colour that is neutral, while a high-volume intersection with poor performance should emit a colour that indicates that something is wrong, such as orange or red.

It can also be seen in the figures that other components of the prototype application were accounted for in the wireframes. These include interactions such as pop-up dialogues, highlighting routes and communicating intersections, as well as a side-panel for showing detailed metrics on intersection and overall city-wide traffic performance. From this stage, the prototype moved to a mid-fidelity phase to further build out the concepts and refine the aesthetic of the application.

While not all of these early-stage concepts were used in the final product, the process of generating, critiquing, and refining multiple ideas proved valuable throughout the year. These initial visualization concepts helped to spawn and guide future implementations, and helped attain a clear understanding of the data being presented. The next step in the design process moved into a medium-fidelity stage, such that navigation and interaction could be explored at a deeper level.

1.2.2 Medium-Fidelity Prototyping

Through the use of prototyping software Balsamiq, a set of interactive mockups were created to both reflect the requirements of the application, as well as ideas and lessons learned from the previous wireframing stage. In this prototype, the main tenets for the final version of the application were created in an illustrative form, to identify key architectural components and the user's interaction within. By doing this in a medium that does not put too much focus on the specifics of visual design, it was easier to concentrate on the core features of the application, and how they satisfy (or dissatisfy) the original requirements.

At this stage, the concept of data layers was explored. Given that there are many varying metrics identified as important to target users (Traffic Engineers), and that these metrics all have a spatial or location-based relevance, the idea of turning on and off layers of data visualization overlaid on a map was critical in effectively demonstrating traffic data. Metrics such as intersection Status, and Flow through an intersection were conceptualized at this stage, with some preliminary UI put together for interacting with these layers. A screenshot of the mockup showing a Status visualization in the Relay app is presented in Figure 2.5.




Figure 2.5: Balsamiq mockup of a data layer in the Relay app.

Additionally, as can be seen at the top of Figure 2.5 the foundations for global app navigation were built out, in the form of a fixed header bar at the top of the screen. This header serves many purposes, including branding for the app, search functionality, as well as tabbed buttons to move between contextually grouped pages. Figure 2.5 shows the active state of the Map button, while the main area of the screen contains a map with relevant data.

Further interaction design was carried out at this stage, in the form of a modular popup window, referred to as the Relay info box. This info box was designed to present deeper metrics for an intersection, and appears when an intersection on the map is clicked. This allows users to quickly inspect any intersection of interest and monitor time-series data in the context of a specific location. Data provided at this stage includes a graph of the flow of traffic through the intersection in each of the cardinal directions, as well as a matching chart that provides predictive data into the near future, based on stochastic methods from the controller. The status of the lights is shown, in addition to numerical car counts and the name of the intersection that has been clicked. A Balsamiq mockup of the info box is presented in Figure 2.6.

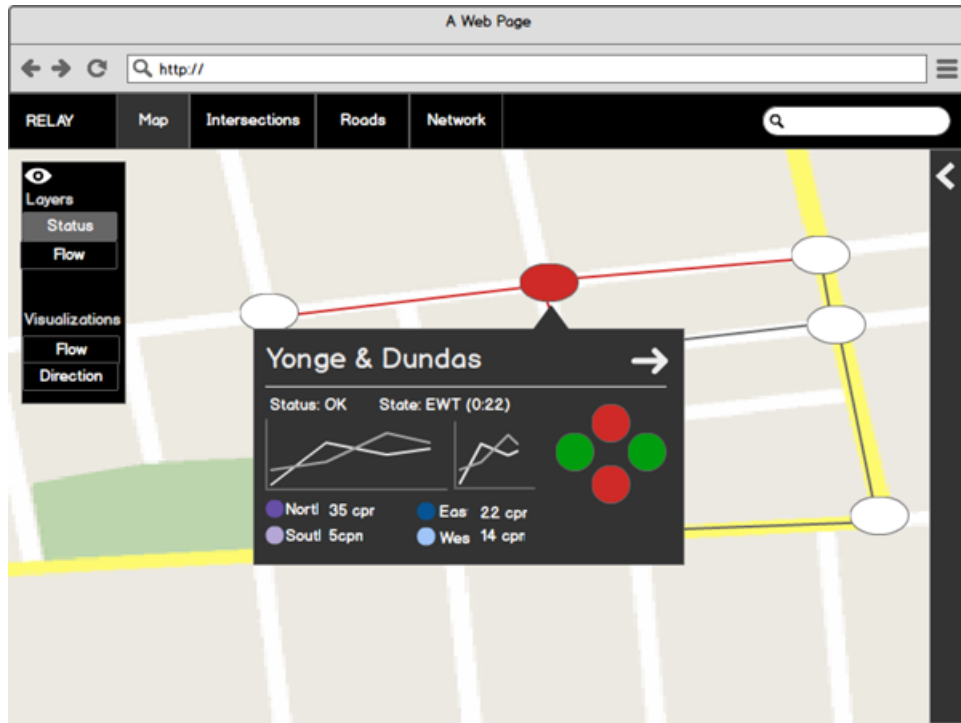


Figure 2.6: Balsamiq mockup of the Relay info box.

While the info box fulfills the need to browse intersections and provide quick snapshots of relevant intersection data, it was important for more advanced users to get an in depth look at this data, such that




Figure 2.7: Balsamiq mockup of the Relay dashboard.

meaningful trends can be detected and acted upon. It was necessary for these users - Traffic Engineers - to view metrics on intersection Location, Name, Type, State, In Flow, Out Flow, Predictions, and more. To accommodate this large quantity of data, a dashboard was prototyped at this stage, setting a preliminary layout for each of the required metrics. This dashboard slides across the screen from the right by pressing the arrow in the info box, and can then be toggled back and forth with the controls in the upper right corner of the screen. A Balsamiq mockup of the dashboard can be found above in Figure 2.7. As various intersections are selected on the map, the dashboard updates it's information to reflect that of the intended location. Time-series graphs update in real-time as data feeds through the system, and an activity log hosts the history of noteworthy incidents at the selected intersection.

With these views on the Map screen prototyped in medium-fidelity, the design progressed into other areas of the application. Moving through the global navigation buttons along the header, an Intersections page was mocked up. The goal of this page is to give Traffic Engineers a holistic view on their network, as well as the ability to filter/search and target individual intersections. This text-based search also allows for insight into any particular intersection by name, as well as the ability to find important groups of intersections or main arteries with many cross streets. This is important because it allows users to locate intersections in two ways: spatially on the map, and now textually in the intersections table. A mockup of

the Intersections page is shown in Figure 2.8, with placeholder data in the intersections table.




Figure 2.8: Balsamiq mockup of the Intersections page.

Finally, the medium-fidelity prototyping stage moved on to build out a layout and basic information for a Network page. The intention of this page is to provide details similar to what would be shown in the info box or dashboard of an intersection, but for the entire network as a whole. This allows Traffic Engineers to view, understand, diagnose, and make decisions for the network on a macro level, in addition to the finer look provided on a per intersection basis. A mockup of the Network page is shown in Figure 2.9.

As can be seen in Figure 2.9, a large activity table is provided on the Network page, to highlight issues that occur at all intersections in the network. These issues include accidents, power outages, abnormal volumes, and more. Additional detail can be seen on the right hand side of the screen, with quantitative network data placed in charts and numerical tables, and colour-coding of network statuses explored at this stage.

Overall, the medium-fidelity prototyping stage was essential in moving the design of the Relay application forward towards its final state. With functional requirements set in place, and now a strong roadmap

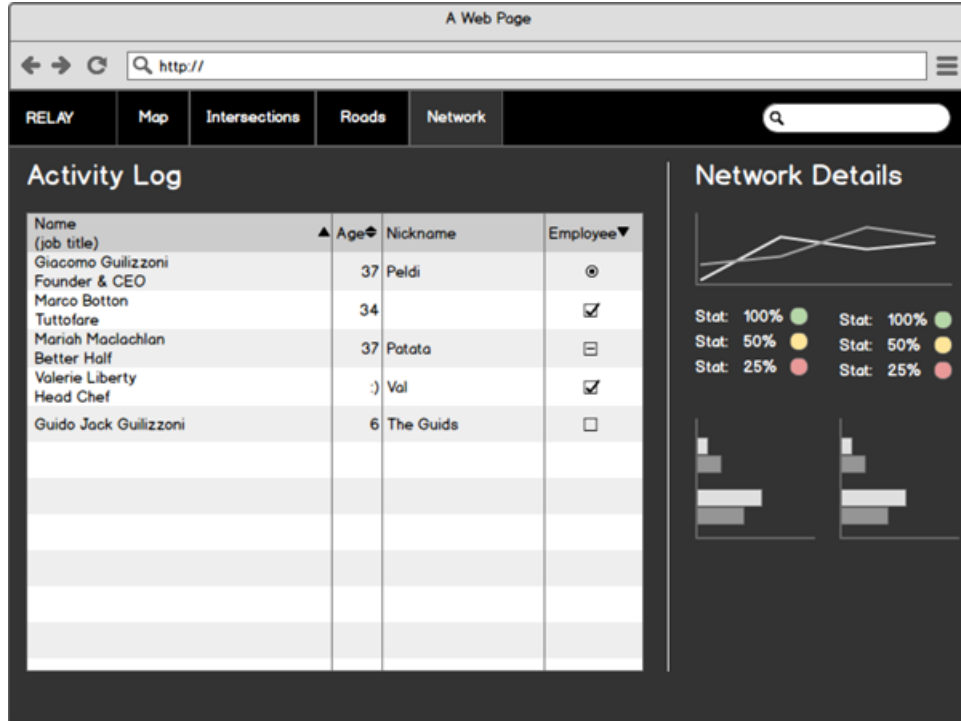


Figure 2.9: Balsamiq mockup of the Network page.

of UI elements, navigation, page layout, and interactions to fulfill the requirements, the application was ready to be built in a high-fidelity form, as described in the following section.

1.2.3 High-Fidelity Prototyping

In preparation for the development of the Relay application, the design of the environment moved into a final high-fidelity prototyping phase. This stage involved a major visual overhaul of the application, focusing on the definition of a unique design language to bring Relay together as a cohesive product. Considerations at this stage include the creation of a custom grid structure, typography, colour, branding, the development of data visualization layers, and repurposing the search functionality in the app. The high-fidelity design was done in Adobe Photoshop, Illustrator, and in the browser using HTML/CSS and JavaScript. These tools allowed for highly accurate creation of assets, development of colour scheme, precise layout of visual elements, and more.

The branding for Relay was created in two main components: the logo, and the typesetting. The logo is a custom glyph used in the application to denote the state of the signal at an intersection. It contains four dots in a diamond pattern, with two fully-coloured dots positioned vertically, and two empty dots with

a medium-weight stroke positioned horizontally. In the application, this glyph would represent the flow of traffic in an East-West manner, while North-South traffic is halted at a red light. This was selected to represent Relay as it both subtly represents what the application does, and provides important information to users in a modern and clean way. The final Relay branding is shown in Figure 2.10.



Figure 2.10: High-fidelity rendering of the Relay branding.

The typesetting was done to represent Relay in a way that would be clear, trustworthy, confident, readable, and modern. The use of custom letter positioning through tracking and kerning methods aided this, while promoting clarity and confidence with the capitalization of each letter in the word. Open Sans, a typeface that is freely available to the public, was utilized for the clean and readable aesthetic it provides. As a sans-serif font set at a semi-bold weight, each letter is clear from any distance, at many sizes, on both screen and in print. This also gives the branding a modern feel, as many serified fonts are often associated with older, more classical and formal scenarios.

The high-fidelity design of the application itself began at a structural level, to set the basis for a consistent, coherent, and unique interface. A grid structure was optimized for a 1440 x 900 pixel resolution screen - a common resolution on modern screens, and the dimensions of the screens used for the Symposium demo. The grid was based upon units of 16 pixels, which was decided to be the lowest common denominator for icon and type sizing to ensure usability from many angles and distances. With a grid structure in place, all following design decisions then had a reference point to follow from. This aided in the purposeful use of negative spacing, symmetry, and balance of each screen composition. The grid on a blank canvas is shown in Figure 2.11.

The first consideration in developing the design language for the application was the typography. Just like the branding, the typography was intended to be clear and readable, but also neutral in tone. Neutrality of the type allows for clear representation of the content itself, which is highly important in such a data-heavy application. As a critical element of the user interface, particular attention was paid to hierarchy




Figure 2.11: 16 pixel grid foundation for the Relay user interface.

of type sizes and weights, all deriving from a smallest line-height of 16 pixels, used in small data labels and body. The typeface Helvetica Neue was used, as it is well-known to provide very clear and readable text at many sizes, with a neutral tone. Additionally, it is a common typeface used in the signage of many major cities around the world. This gives the application a recognizable and associative feeling, which is especially important for an application that deals with traffic conditions in urban areas. An example of an intersection name set in 24pt Helvetica Neue Bold is shown in Figure 2.12.



Figure 2.12: An intersection name set in 24pt Helvetica Neue Bold.

The high-fidelity prototyping process continued with the development of colour-scheme for the application. Given the importance of visualizing data in the Relay app, it was decided that all non-data elements of the UI - the map, modular overlays, and header, for example - should be subdued and clearly grouped together as functionally similar elements. To accomplish this, a black and white UI was established for all control elements, providing a clean environment for navigation, while vibrant, highly-saturated colours were used for charts and data layers. This technique effectively created a high-contrast environment where areas that require visual focus, attention, and thought - i.e. the data - were the most salient elements on the screen, and any layout or navigational components were done in a clear black and white manner, similar to signage found on the streets of many major cities. Therefore, all contextually similar elements were easily

grouped and identifiable through their similarities or differences in colour. An example of a chart showing the flow of traffic through an intersection over time is illustrated in Figure 2.13. Notice the vibrancy of the coloured data, and its contrast with the structural page elements surrounding it.




Figure 2.13: A chart showing the flow of traffic through an intersection over time.

The data visualization layers were next to be developed at the high-fidelity stage. These visualizations took the form of heat maps, as well as quantitative, relative, and boolean glyphs, representing data on the Status, Flow, and Queue Length at a given intersection. The first visualization layer, showing the Status of intersections, used small dot glyphs of different colour to display the boolean status of a light. By overlaying a dot at each intersection, with a corresponding white or red colour to denote proper working condition vs. an error, such as a power outage, an easily digestible visualization was created. This empowers Traffic Engineers to easily monitor their network, and to detect any anomalies that may have occurred. A screenshot of the Status layer is shown in Figure 2.14.

The second data layer to be designed was a heat map representing the flow of traffic through intersections. This was accomplished on a per-intersection basis, and was intended to provide a macro view of the relative traffic density across areas of the city. This view is most useful for those looking for an overview of general traffic conditions, such as a browsing Traffic Engineer, or a consumer looking to plan a route through the city. A screenshot of the first Flow visualization layer is shown in Figure 2.15.

Following this Flow visualization, a second implementation of traffic flow was developed, this time looking at flow between intersections, rather than at or through any particular intersection. This was accomplished by taking predictive data from the controller, and approximating the position of cars along roadways given the knowledge of their entry and exit directions from each intersection. Each car is plotted as a low-opacity glyph of concentric circles, with a bright green core, a lighter green exterior, and a soft blue

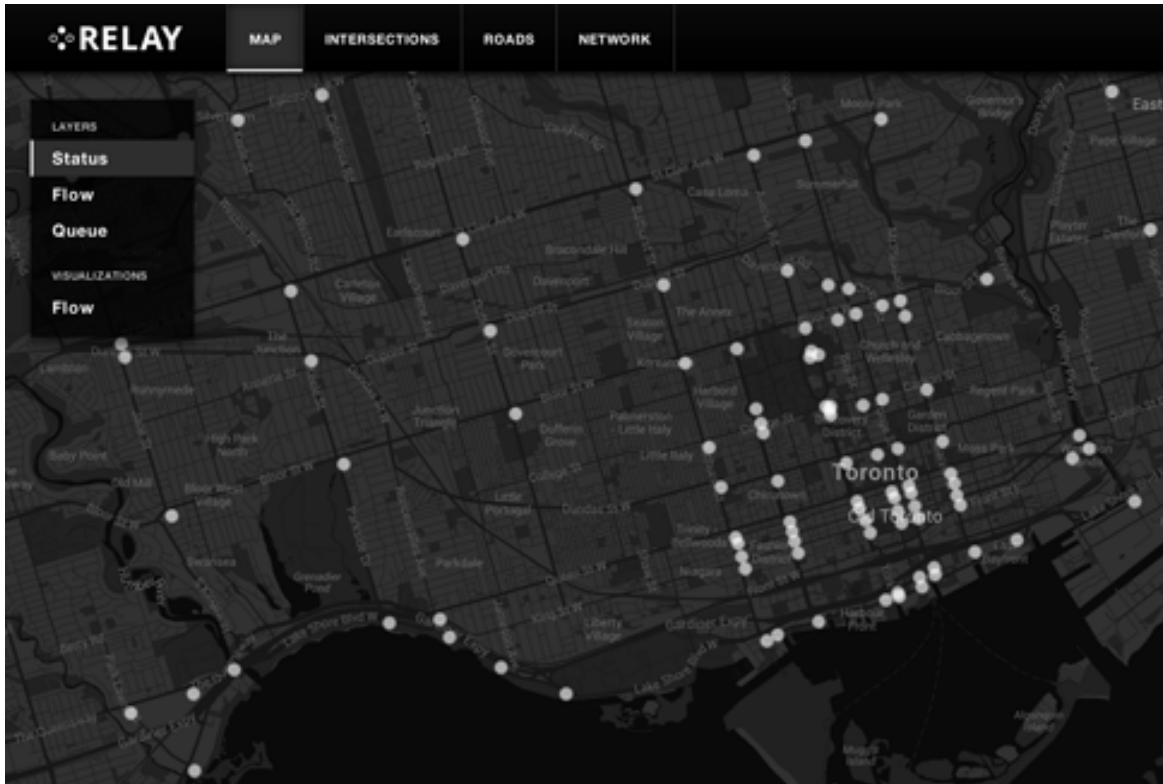


Figure 2.14: A screenshot of the intersection Status data layer.



Figure 2.15: A screenshot of the flow per-intersection heat map.

outline, denoting the estimated nature of the position of any given car. Given the large number of vehicles to be accounted for, a low opacity was selected for each individual glyph, so that as they overlap with one another, the intensity of the colour increases so as to illustrate the increasing density of traffic in an area. A screenshot of the Flow between intersections data layer is shown in Figure 2.16.

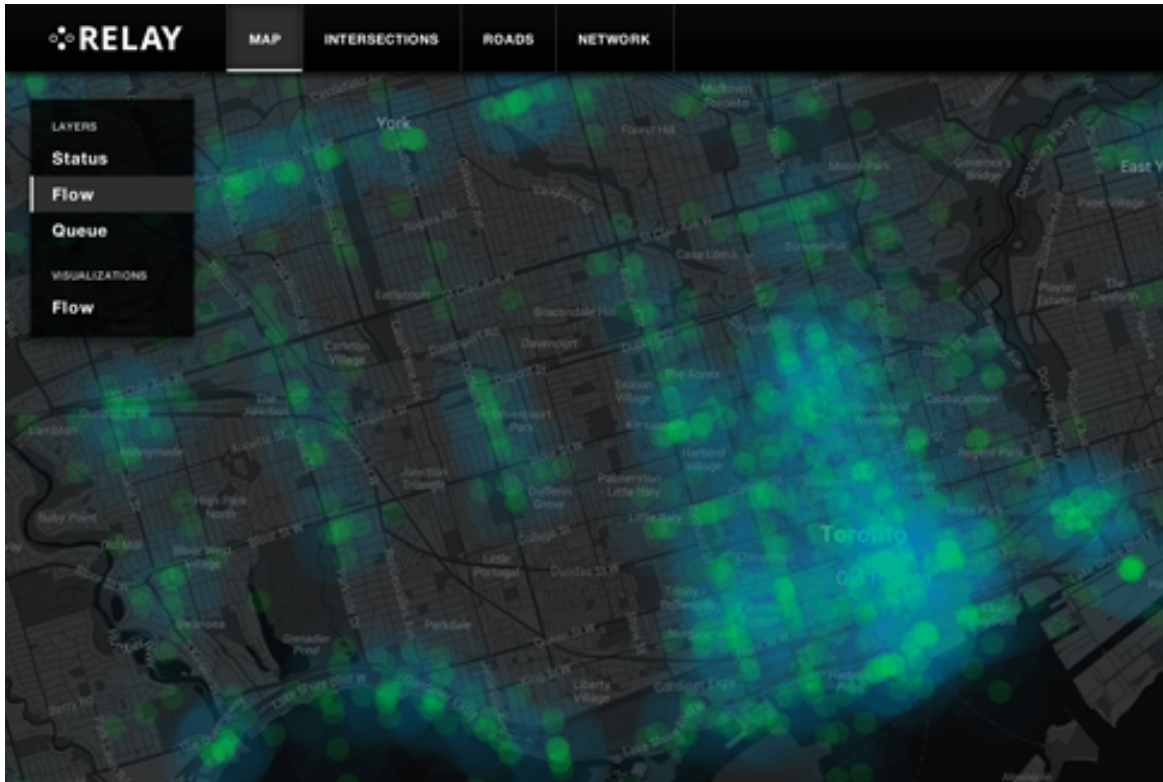


Figure 2.16: A screenshot of the flow between intersections heat map.

As can be seen in Figure 2.16, the natural visualization of major roadways and arteries through the cities is evident, which is an important result of this visualization. A wide range of traffic patterns can be derived from this, thanks to the high-resolution of data points that give a continuous feel to the data, as opposed to the limited two-stage discrete options currently available through applications such as Google and Apple Maps.

Finally, the final visualization layer looked to showcase the data representing Queue Length or average wait time at an intersection. The goal of this visualization was to produce glyphs that could show both quantitative measures, as well as give a sense of relative weighting between intersections. A four-pronged glyph was developed, with three short prongs extending in the East, South, and West directions, while a variably sized prong extends to the North. The length of the North prong corresponds to the average wait time at a given intersection, while the three shorter prongs are of equal length and serve as a stand for the

quantitative North measure. This creates a pseudo 3-dimensional effect, and gives the user a unique look at the changing landscape of queue lengths in the city - similar to how a city skyline may look with many tall buildings. A screenshot of the Queue Length visualization layer is shown in Figure 2.17.

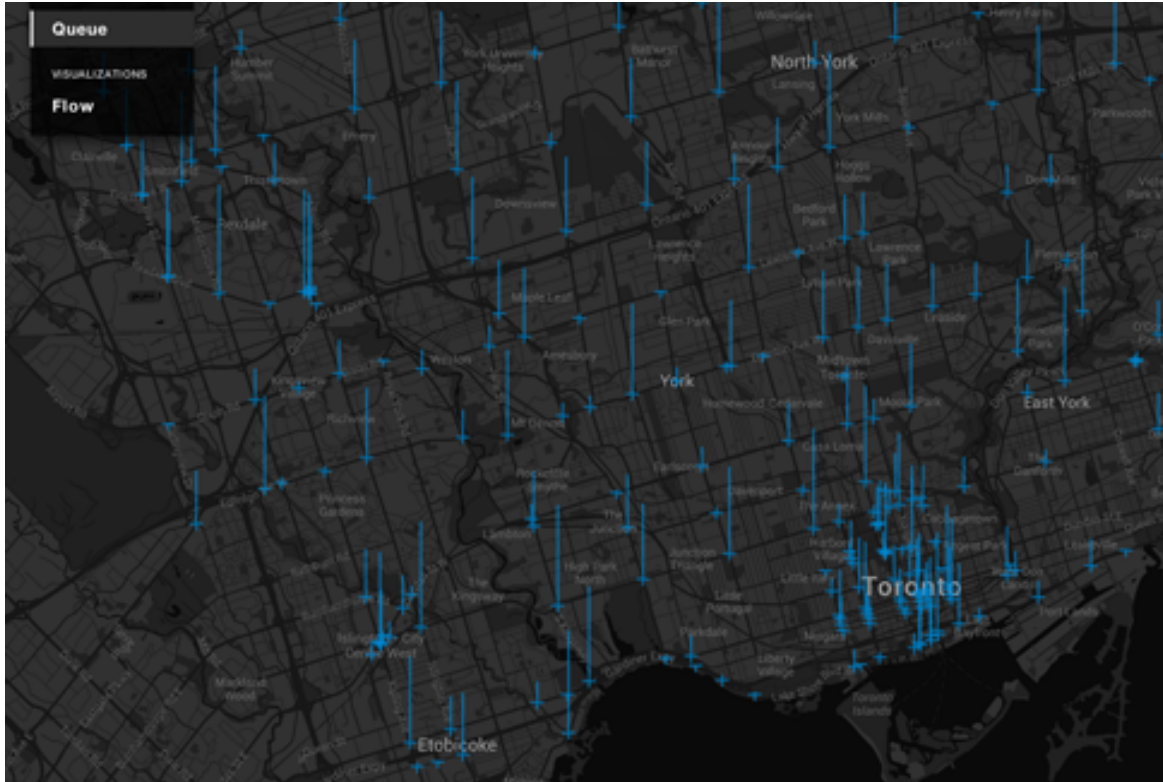


Figure 2.17: A screenshot of the average queue length at each intersection.

Following the design of the data visualization layers, the high-fidelity stage looked at the search functionality of the application, and focused attention on repurposing its form to better reflect its function. Previous prototypes held the search bar in the header across all pages of the application. This proved to be troublesome, as some pages did not require any search functionality, while the ones that did actually served as text-based filters within the context of the current page. The high-fidelity design saw the removal of the search bar from the global navigation header, and into the upper section of the Intersections and Network pages. This way, the text-based tables found on these pages could be easily searched/filtered, without the confusion of entering text into the global header, where a new page would be the expected result.

An additional feature that made its way through the high-fidelity phase is the Connection Status window. This window is activated through the gear icon in the top right corner of the screen, and is useful for monitoring the status of the connection between the front-end and back-end of the application. This window is positioned in the top-right of the screen so as to avoid covering any important visual elements, and

is padded in from the top and right sides of the screen to maintain access to the dashboard toggle button. The connection status window identifies where the data is coming from: either through a Web Socket or through the application's REST API. It also identifies time since the last data update, both locally on a selected intersection, and globally over the entire network. A screenshot of the Connection Status window is shown in Figure 2.18.

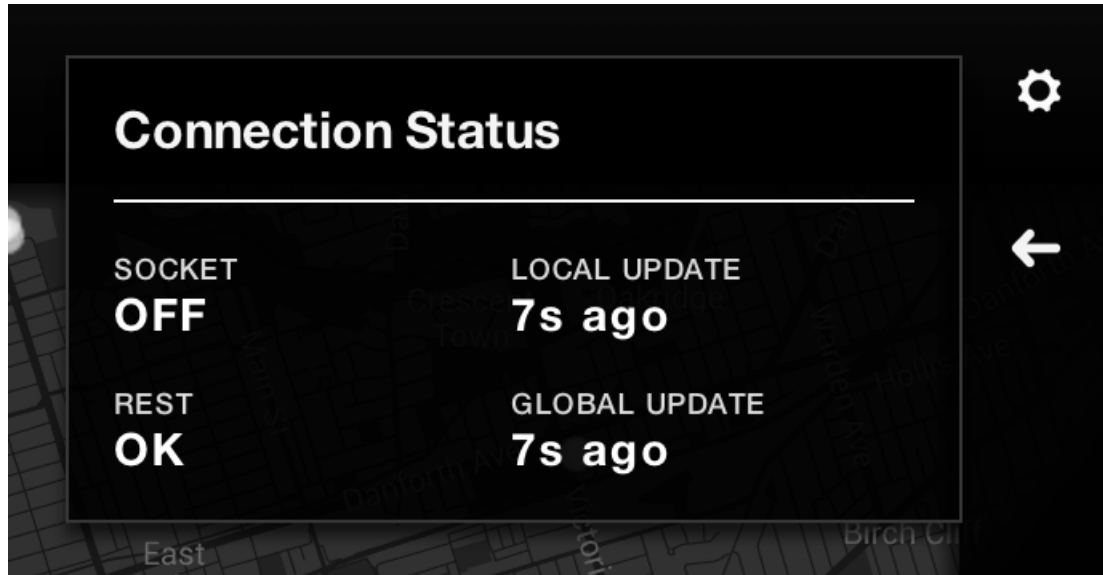


Figure 2.18: A screenshot of the connection status window.

Lastly, the intersection dashboard was refined at a high-fidelity stage. The specific metrics contained in the dashboard were iterated on and refined from the medium-fidelity stage to accurately reflect the data coming from the back-end controller. This involved splitting the Flow metric into In-Flow and Out-Flow at an intersection, showing the current and next intended State of the traffic lights, and providing a new Turn Predictions matrix, which gives the predicted quantity of cars moving through each permutation of turns based on stochastic data from the back-end. In-line with the medium-fidelity prototype, the Activity Log is provided giving consideration to readability and interaction, with subtle hover states on individual table rows to improve salience of desired fields. Additional information includes the Intersection ID and the Intersection Type, two identifiers that are often important to Traffic Engineers. A screenshot of the high-fidelity dashboard is shown in Figure 2.19.

With these elements in place, the final developmental/functional touches to the Relay application were ready to be built for the demo, as showcased on Symposium day. As a unit, the application is held together using the modular design elements described above, which makes the application easily scalable

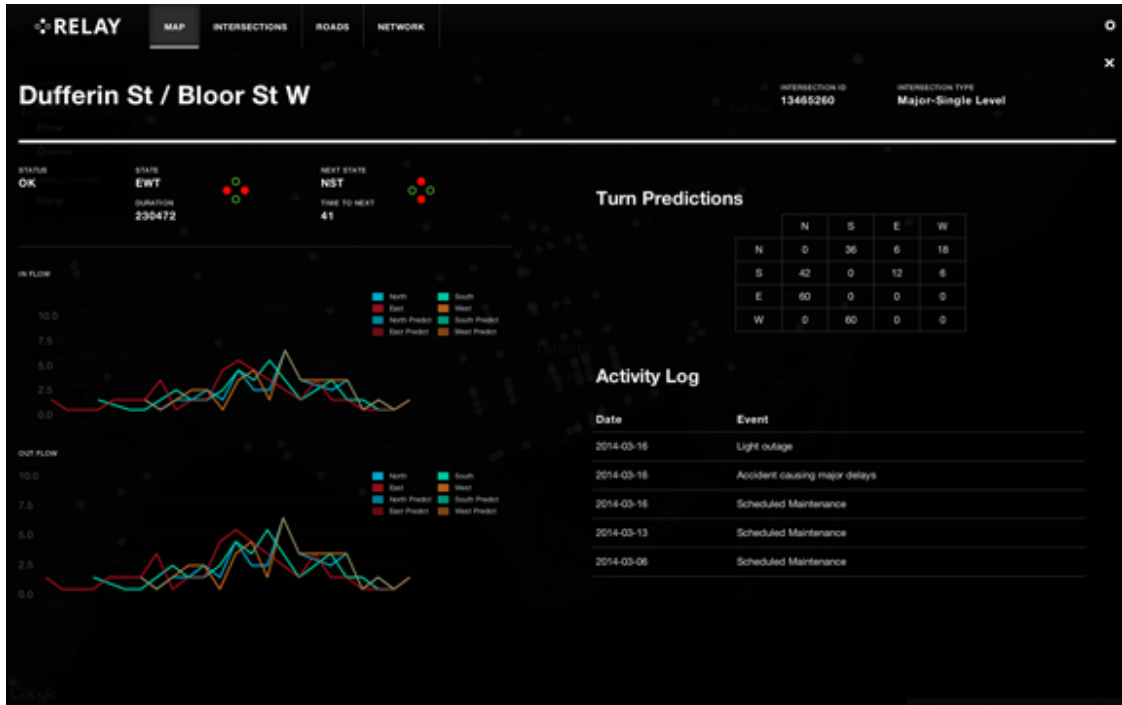


Figure 2.19: A screenshot of the Relay intersection dashboard.

and modifiable for individual implementations or changing metrics and requirements. With a unique and well-defined design language established, further extrapolations of this implementation have a library of best practices and UI elements that can be carried forward into new scenarios.

1.3 Application Architecture

Application interfaces for critical infrastructure must be highly robust and reliable. Furthermore, it is important the the systems are able to scale in order to accomodate the inevitable growth of the system. These main ideas drove the design of the front-end (web application) application architecture.

Given the breadth of choices in software technology, it was necessary to use a quantifiable selection process for determining the appropriate solutions from a group of candidates. The three main areas of the front-end application which required such a process are the Framework, Mapping Library, and Visualization Library. The selection process for these is discussed in the following sections.

1.3.1 Framework Selection

Software frameworks are code libraries which provide functionality that enables developers to more easily implement a specific software pattern. One of the most common software patterns is the Model-View-Controller (MVC) pattern, which is commonly used in complex and large software applications to ensure the application is reliable, structured, and can scale effectively. Similarly, front-end software frameworks allows for the development of structured and scalable applications by providing software pattern utilities. Given the critical nature of our application, it was determined that the use of a framework would be appropriate for ensuring the scalability and reliability of the application.

Four JavaScript frameworks were selected for comparison. Javascript MVC is a mature, well-known front-end MVC implementation which is build off of jQuery, a JavaScript Utility Library, which offers a well-rounded feature base. DOJO is another mature and well-known framework which does not offer an MVC-style framework, yet is class-oriented. Backbone.JS is a relatively new Model-View (MV) framework built off of the Underscore.js utility library. Finally, Angular.js is a new and highly functional front-end framework which loosely follows an MVC pattern.

The frameworks are scored across five weighted features. The scores are justified and a winning framework is selected.

1. **Size** - The use of large JavaScript libraries can cause performance issues within the application. The selected library should be as small as possible. Candidate frameworks will be assessed based on the size of their encompassing code files. This feature has been given a weight of 0.5.
2. **Independance** - The selected framework should be self-contained and not require many prerequisite JavaScript Libraries in order to function. The presence of many prerequisite libraries increases the size of the application and will slow down performance. Candidate frameworks are assessed based on the number of prerequisite frameworks and plug-ins required for the framework's application within the proposed solution. This feature has been given a weight of 0.5.
3. **Functionality** - The selected JavaScript framework should provide a rich set functionality which is not available within the native Javascript language. Important functionality includes utility functions, templating, Document Object Model (DOM) binding, client-server communication methods, and data structure classes. Candidate frameworks are assessed based on the both the quantity and quality of the functionality provided, in the context of the proposed solution. This feature has been given a

	Angular.js	Backbone.js	DOJO	JavaScript MVC
Size (0.5)	5 (2.5)	3 (1.5)	2 (1)	2 (1)
Independence (0.5)	5 (2.5)	4 (2)	3 (1.5)	4 (2)
Functionality (0.8)	3 (2.4)	5 (4)	4 (3.6)	4 (3.6)
Modularity (0.8)	2 (1.6)	4 (3.6)	4 (3.6)	4 (3.6)
Implementation Overhead (0.7)	4 (2.8)	4 (2.8)	3 (2.1)	4 (2.8)
Totals	11.8	13.9	11.8	13.0

Table 2.1: Weighted Decision Matrix for JavaScript Framework Selection

weight of 0.8, as it is important that the framework provide sufficient functionality to implement the solution.

4. **Modularity** - It is important that the selected JavaScript framework supports the creation of modular code, given the complexity of the proposed solution. The selected MVC framework should encourage an established code pattern, such as and MV- or MVC-style structure. Candidate frameworks are assessed based on their similarity to known software architecture patterns, such as MV or MVC. Due to the importance of scalability and modularity in the proposed solution, this feature has been given a weighting of 0.8.
5. **Implementation Overhead** - The selected JavaScript framework should not require excessive amounts of overhead code in order to successfully suit the needs of the proposed solution. The selected framework should have as little overhead as possible. Candidate frameworks are assessed by examining tutorial code implementations for each framework, to determine the proportion of development required as overhead compared to application-specific code. This feature has been given a weight of 0.5.

The frameworks are presented in a weighted decision matrix below. Each attribute is weighted out of 5 points for each of the candidate frameworks. Weightings for each attribute are on a scale from 0 to 1.

Angular.js is the smallest candidate framework at 36kb. Angular.js also has no dependency JavaScript libraries, and is thus completely independent. Although Angular.js offers rich DOM and interaction functionality, it provides few utility features for data structures and architectures. Angular.js does not follow a rigid MVC or MV structure, which was a major downfall for the framework. Angular.js does not require much overhead code to implement, primarily because it does not follow a rigid software architecture, and thus is more functionally-oriented.

The Backbone JavaScript framework is 64kb in size. Backbone.js relies on Underscore.js and jQuery.js two relatively large JavaScript utility libraries. However, both of these libraries will be used in the application regardless of the choice of framework, thus are not direct detriment of this candidate. Backbone.js follows

a MV framework style, where both controller and view functionality are both kept in the view. This pattern is a good compromise between no framework and an extremely heavy MVC structure. Backbone.js's two utility libraries provide rich functionality for working with the DOM model, and structured objects. Furthermore, Backbone provides a native collection-oriented pattern developed specifically for handling large numbers of similar objects, which fits well with the applications need to support many intersection and road models. Finally, Backbone.js has minimal development overhead considering it's MV structure, as Backbone.js provides a range of default functionality for it's classes.

The DOJO JavaScript framework is 155kb in size, making it the largest of the candidate frameworks. DOJO requires no prerequisite JavaScript Libraries, yet many of DOJO's features are in add-on libraries. This reduces the effectiveness of DOJO's core code, yet means a wider range of features are available as the feature set distributed and open. However, DOJO's functionality is not aimed specifically at dealing with large collections of data, and does not offer strong support for client-server synchronicity. In terms of modularity, the DOJO framework supports the creation of class objects yet does not explicitly support an MV- or MVC-style framework. DOJO requires a significant amount of implementation overhead, which makes it suitable for extremely large and diverse applications. However, this level of overhead is less suited to the complex yet relatively consistent nature of our application with respect to object and data needs.

JavaScript MVC is 105kb in size. JavaScript MVC relies on the jQuery.js JavaScript library, however this library is already a requirement for the proposed solution. JavaScript MVC is composed of four sub-frameworks, each of which specializes in a different set of features, and together provide sufficient functionality for the proposed solution. However, like DOJO, JavaScript MVC lacks the specific functionality for dealing with large amounts of similar objects. As suggested by the name, JavaScript MVC follows a strong MVC model in it's implementation.

As seen in table 2.1, Backbone.js is the winning candidate. Backbone is not the smallest candidate, yet it's feature set and MV framework fit well with the needs of the proposed solution. Specifically, Backbone's use of already-implemented JavaScript utility libraries to provide enhanced data structure tools and algorithms, as well as it's specialized collection template for dealing with large amounts of data objects. Furthermore, Backbone.js provides functionality for client-server communication for said collection structures. Finally, the overhead associated with Backbone.js is acceptable given the strong fit of the functionality with that of the proposed solution.

1.3.2 Mapping Library Selection

Next, a front-end mapping library needed to be selected. The proposed solution required that information be viewed in a geographic context. Thus, it was necessary to select a client-side library which would support these needs. Generally, it was important the the library be highly-customizable, as the data presentation layers would be completely customized. Furthermore, the library had to be high-performance to support the needs of our application. Finally, it was necessary to find a map hosting service which could support the library, as setting up a hosting service was not in the scope of this solution.

1. **Style Customizability** - The selected mapping candidate should be highly customizable. Colour and styling for map features such as roads, buildings, water, and land must be allowed. This feature has been given a weight of 0.5, as it is not crucial yet still important to the aesthetic value of the interface.
2. **Feature Customizability** - The selected mapping candidate must allow for the customization of built-in features, such as markers and info boxes. Markers should be completely customizable, through the use of either a scalable vector graphic (SVG) or HTML and CSS. Info boxes should be customizable using HTML, CSS, and JavaScript. This feature has been given a weight of 0.7 for it's importance in defining the customizability of crucial interface features.
3. **Layer Customizability** - The selected mapping candidate must allow for the customization of map data layers. The candidate must support multiple simultaneous layers of data. It should also support the display of both point-based and vector-based data. This feature has been given a weight of 0.9, as the value of the application is directly related to the presentation of the various data layers.
4. **Tile Hosting Support** - The selected mapping candidate must either support tile hosting as part of its service, or be compatible with an existing tile hosting service. Ideally, the selected candidate offers both tile hosting and serving in the same package for convenience. This feature has been given a weight of 0.6, as it is not a crucial feature yet saves significant effort if present.
5. **Price** - The selected mapping candidate must be cost effective. The service should be free, or minimize the cost of tile serving and hosting. This feature has been given a weight of 1 due to the highly restrictive budget for the proposed solution.

Four Mapping Libraries were considered as potential candidates: Google Maps API, Leaflet.js, Map-Box, and TileMill in conjunction with CloudMade.

	Google Maps API	Leaflet.js & CloudMade	MapBox & TileMill
Style Customizability (0.5)	3 (1.5)	3 (1.5)	5 (2.5)
Feature Customizability (0.7)	4 (2.8)	5 (3.5)	3 (2.1)
Layer Customizability (0.9)	5 (4.5)	4 (3.6)	3 (2.7)
Tile Hosting Support (0.6)	5 (3.0)	3 (1.8)	3 (1.8)
Price (1)	5 (5.0)	4 (4.0)	3 (3.0)
Total	16.8	14.4	12.1

Table 2.2: Weighted Decision Matrix for Mapping Library Selection

The Google Maps API is the industry standard of mapping libraries. It is the same framework which runs Google Maps, and is thus well-supported and reliable. The Google Map API allows for styling customizations on-the-go through its JavaScript API. Unfortunately, not all aspects of the map are styleable. The API provides functionality for both markers and info boxes. The Google Maps API allows for custom CSV marker icons, however the info box is not easily customizable. However, Google Maps does offer a plug-in info box, which is fully customizable. Tiles are hosted directly through Google Maps, requiring no extra work for tile hosting support. Finally, the Google Maps API is a free service with very high courtesy limits.

Leaflet.js is a map-oriented JavaScript Library which provides a variety of map-related features on top of a tile service. Leaflet.js does not host map tiles, and thus it must be used in conjunction with a tile hosting service such as CloudMade. As such, styling the map must be done through a CloudMade's third party service which is restrictive in styling. Leaflet.js offers rich map functionality, including markers and info boxes. Furthermore, leaflet.js' map features are highly-customizable, as it is an open-source project. However, Leaflet.js is not geared towards handling data layers. Leaflet.js is a free library.

Mapbox is a new JavaScript mapping library which works in conjunction with TileMill, a tile styling service. TileMill is designed specifically for creating highly-customized map styles, and thus offers very rich styling capability. Due to MapBox's immaturity as a platform, its feature set is limited and not very customizable. The MapBox JavaScript Library is not structured towards handling data layers, rather individual data points. Furthermore, TileMill tiles must be hosted using a third part platform, or through MapBox directly, which required a paid membership for the service.

The candidate mapping libraries are compared in Table 2.2. The Google Maps API is the winning mapping library. Although the Google Maps Library lacks full native customizability for both its map styling and map features, sufficient customizability is found within the native functionality and through the use of add-on libraries. The Google Maps API provides excellent data layer functionality, which is unparalleled in feature breadth and support. Finally, the Google Maps combined tile hosting and JavaScript functionality

provided within a free service make it the clear choice.

1.3.3 Visualization Library Selection

Given the proposed solution's heavy focus on data presentation and visualization, it was important to select a visualization library which supported a wide range of data visualization formats. It was necessary to select a library which allowed for heavy style customization, including line format, colours, grid styles, and legend displays. Considering that traffic information exists in the temporal space, it was important that the library handle time series data well. Specifically, functionality for handling timestamp formatting and real-time or high-frequency updating.

The candidate visualization libraries are assessed across 5 feature categories.

1. **Prerequisites and Add-Ons** - The selected visualization library should require minimal additional plug-ins or libraries to run. All of the necessary functionality should be available within the core library files. Candidate libraries will be assessed based on the number of additional files or libraries required to perform. This feature has been given a weight of 0.5, as the functionality of the library is relatively more important than the quantity and size of the library's dependencies.
2. **Implementation Complexity** - The selected visualization should minimize the amount of code required to produce the necessary visualization. Unnecessarily long setup scripts may cause performance issues in the code. Candidate mapping libraries will be assessed by examining example code and tutorials in order to understand the complexity of visualization implementation. This feature has been given a weight of 0.8, due to the frequency and extent of visualization use within the proposed solution.
3. **Time-Series Support** - The selected visualization library should have rich time-series support. The library should support timestamps in various data formats and allow for the display of time values in a variety of formats. Candidate mapping libraries will be assessed based on the extent of the documented time-series formatting functionality available for the library. This feature has been given a weight of 0.9, as time-series support is crucial for the proper display of information within the proposed solution.
4. **Style Customizability** - The selected visualization library must allow for highly customized visualizations. The library should allow for the customization of visualization style, line style, legend style and placement, grid style and placement, and axis style and placement. Candidate mapping libraries will be assessed based on the extent of documented customizable features available for the library. This

feature has been given a weight of 0.8, due to the importance of data presentation style in the success of the proposed solution.

5. **Ease of Updating** - The selected visualization library must support real-time updating of the data within the visualization. The library should either explicitly support data updating through JavaScript methods, or allow for the replacement of the data object and the redrawing of the visualization. Candidate libraries will be assessed based on the perceived presence of such functionality within the library documentation. This feature has been given a weight of 0.9, as the proposed solution relies on the real-time display of information.

Four JavaScript-based visualization libraries were assessed as candidates.

Chart.js is an HTML5-based visualization library, focused on delivering implementations of several core data presentation formats, e.g. bar charts and line charts. Chart.js is a stand-alone library, requiring no outside code. The implementation of visualizations in Chart.js is easy and requires minimal amounts of code. Chart.js does not support any formatting for axis labels, including support for date formatting and time-series support. Chart.js provides an intermediate amount of formatting options for visualizations, however the restriction of the library to 6 main visualization formats limits the presentations options while using this library. Furthermore, Chart.js does not support the updating of chart data in any way, and would require the chart to be recreated in order to display new data.

D3, or Data-Driven Documents, is the most well-known and feature-rich HTML visualization library today. D3 provides an extremely wide range of visualization formats which are completely customizable. D3 has no pre-requisite libraries, and is a stand-alone library. D3's main downfall is the complexity of implementation, as basic data presentations can require up to 300 lines of code. As an extension of D3's extreme customizability, it is possible to use time-series data, however the library provides little native formatting support. D3 provides the widest range of visualization options by far. Finally, D3's support of data updating depends on the visualization method used, however it is generally poor due to the static nature of the library.

Flot is a JavaScript-based visualization library built on top of the jQuery utility library. Flot is a medium-complexity visualization library which emphasizes informative and real-time visualizations. As mentioned, the Flot library requires jQuery to run, however jQuery is already a prerequisite for multiple other libraries and feature in the proposed solution. Flot charts are not designed to be extremely easy to implement; their implementation complexity is significant, yet more efficient than other candidates and understandable

	Chart.js	D3	Flot	Rickshaw
Prerequisites & Add-Ons (0.5)	5 (2.5)	5 (2.5)	4 (2.0)	4 (2.0)
Implementation Complexity (0.8)	4 (3.6)	1 (0.8)	3 (2.4)	5 (4.0)
Time-Series Support (0.9)	1 (0.9)	3 (2.7)	5 (4.5)	3 (2.7)
Style Customizability (0.8)	2 (1.6)	5 (4.0)	4 (3.6)	3 (2.4)
Ease of Updating (0.9)	1 (0.9)	2 (1.4)	4 (3.6)	1 (0.9)
Total	9.5	11.4	16.1	12.0

Table 2.3: Weighted Decision Matrix for Visualization Library Selection

given the functionality available. Flot natively offers moderate time-series formatting support: Flot accepts JavaScript timestamps as time-series data, and is able to format said data in a variety of styles. With the addition of a jQuery plugin, the range of styles is greatly increased. Flot provides a moderate amount of customizability for their data presentations. However, due to the open-source nature of Flot, the inner styling variables are publicly known, which allows them to be customized far beyond what is supported through public JavaScript methods. Finally, Flot handles data updating very well, with explicit JavaScript functions for updating data objects and re-rendering visualizations.

Rickshaw is a new, static visualization library build using D3 functionality. It aims to offer less complex implementations of basic data visualizations using the D3 framework. Obviously, D3 must be loaded in order for Rickshaw to run, which brings the negative performance implications of loading such a large library as D3. Rickshaw data presentations are exceptionally easy to define and require minimal code. However, the immaturity of the library means that only a few plotting styles are supported, greatly limiting the library’s capability. Furthermore, Rickshaw is a private library and is thus less accessible to customizing beyond the limited exposed styling options. Rickshaw does provide time-series formatting support, however it is limited in scope and buggy as the library is not yet a mature project. Finally, due to the static nature of the underlying D3 platform, the Rickshaw library has no support for updating the data object, and thus would require the recreation of the visualization.

Assessment of the candidate visualization libraries is done through a Weighted Decision Matrix, as seen in Table 2.3. As seen in Table 2.3, Flot is the winning visualization candidate. Flot’s exceptional functionality with regards to time-series data and data updating fit well with the needs of the proposed solution. Furthermore, the Flot visualization library provided a suitable balance of complexity and customizability for the proposed solution.

1.3.4 Architecture Development

The architecture for the front-end application saw three major iterations. Each phase represents a more complex yet refined application architecture. Tables 2.20, 2.21, and 2.22 document the Javascript architecture at each iteration. This section only discusses the progression of the JavaScript class architecture, and does not discuss the HTML or CSS development which happened in parallel. However, all other front-end development was done to support the JavaScript class structure, and thus discussion of their implementation is not necessary.

Due to the use of Backbone.js as the JavaScript framework, the application follows an MV architecture, where only models and views exist as distinct classes, and controller functionality is included in the view object. Backbone.js includes a specific model pattern called a collection, which is an object responsible for handling multiple models of the same class. The solution leverages the functionality provided by the collection pattern, and they have been appropriately categorized in the architecture diagrams. However, collections are similar to models in the sense that they have no controlling functionality, and thus are treated like models in the sense that each has an accompanying view class to handle its controlling functionality.

The 'User Entry' mark within the class diagrams indicates the entry point to the initialization of the JavaScript code. This is done using a call to the script (iteration 1) or through the initialization of the indicated class object (iterations 2 and 3), which then progressively initializes all other classes within the structure. The arrow direction indicates the direction of ownership within the classes. That is, class A pointing to class B indicates that class A owned one or multiple instances of Class B. At the bottom of each diagram, database connection points are indicated. The presence of a line between a model and a database indicates the ability for the model to query the database. Finally, views which are directly owned by a collection view are stacked. This is to show the collection view's 1 to n relationship with the view classes in the implementation.

The first major implementation was a proof-of-concept interface developed for the Fall Demonstration and Presentation. The goal was a low-fidelity simulation of the core functionality in the solution. Thus, the code was developed in a less-structured manner so it could be changed quickly as various implementation strategies were tested. As seen in Figure 2.20, all functionality was developed in a single, classless script. No specific class pattern was implemented at this stage.

The second major implementation was completed for the Winter Demonstration and Presentation. A more mature class architecture had been developed, seen in Figure 2.21 as many of the necessary utility




Figure 2.20: An architectural diagram for Iteration 1.

libraries had been selected. Furthermore, Backbone.js had been selected as the JavaScript framework of choice, and significant effort had been made into transitioning the original functionality into the Backbone.js pattern.

The App View class is primarily responsible for handling the initialization of the application, and acts as a wrapper for all other functionality. This iteration contained a single screen, and the screen was appropriately given a View class to manage its interface and interactions with the model. Furthermore, intersections pulled from the database were stored in a Collection class, where each intersection was represented with an intersection model. The intersection class created an intersection view for each model, which handled direct interactions from the interface on a specific intersection.

It is important to note that the model and collection classes are initialized at the app level. This allows for the model objects to be used as singleton instances of the network information, removing the need for redundant data in the interface, as well as reduction in data transfer between the client and server. Furthermore, this strategy improves the robustness of the client application though restricting server communication to a single outlet.

The third and final iteration of the architecture, as seen in figure 2.22, was developed for the final product demonstration and symposium. This architecture represents the full implementation as outlined and discussed within this report.

The overall structure of the application is similar to that in the second iteration. The App View

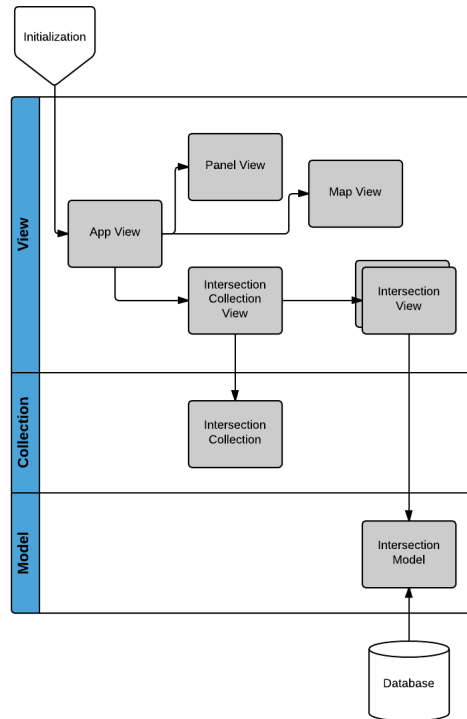


Figure 2.21: An architectural diagram for Iteration 2.

class is responsible for the high-level initialization and management of the application instance. The App View also owns both the main model collections and page views. The Road Model, Road Collection, and associated views were added to support the display of road information within the interface. Furthermore, three new pages were added in addition to the original map page view: Intersections Page View, Roads Page View, and Network Page View. These pages utilize the existing data, yet present fundamentally different views and functionality around the data, supporting many of the non-crucial functional requirements. The addition of non-page and non-model views, such as the Panel View and the Activity View, encourage the application of Backbone's MV framework for the sake of uniform code patterning.

The value of storing one instance of the intersection and road collection objects is now clear, as the four pages now feed off of the same two data representations. This also means that all updating and synchronization of information between the client and server has been centralized, and thus the data is guaranteed to be consistent across the application.

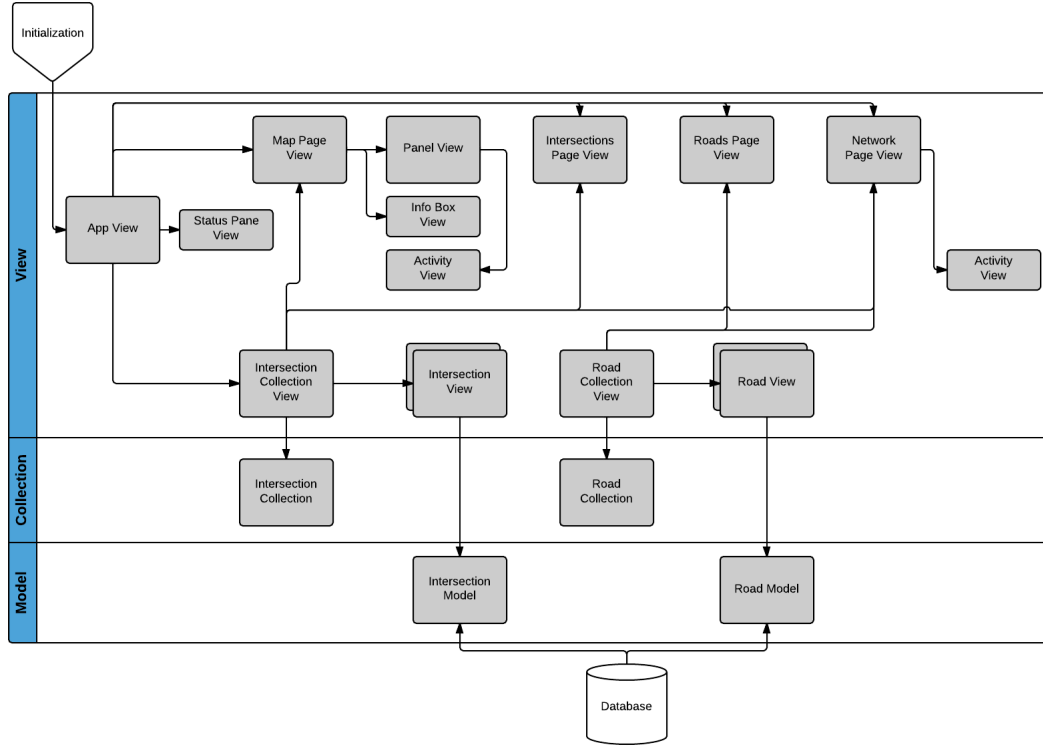


Figure 2.22: An architectural diagram for Iteration 3.

2.0 Back End

3.0 Deep End

3.1 Learning in Relay

3.1.1 Design Parameter Updates

Each agent within the system has multiple design parameters that will constantly update. With each cycle – behaviour change – predictions and estimations are made to recalculate these parameters. The following sections explain how these variables are learned and updated.

3.1.2 Time Delay Estimation

The first important parameter needed to calculate predictions is the time delay estimate between two agents. This can be thought of as the estimated time to travel from one intersection to a neighbouring one. However, this is not a fixed time, so the delay is modelled using a gamma distribution, as discussed previously.

The first step in performing this estimation, is determining which portion of the downstream signal is relevant meaning what part of the signal should be considered for the estimation. This is required because it must match the departing signal from the upstream intersection. To extract the downstream signal, an upstream sub-signal is first selected. It was determined that the most efficient way to do this is to select the upstream departure signal created by the previous behaviour cycle. To ensure that events are not missed, the downstream signal is determined by time-shifting the upstream signal by the expected value of the previous time delay distribution, and a small buffer is added to each end.

After the two signals are extracted, the estimation process can be performed. Equation [time-delay-opt] outlines the optimization problem, using least squares. The goal of this optimization is to convolve the departure signal with a gamma distribution to best estimate the downstream arrival signal – to minimize the difference between the real signal and estimated. The design parameters are: shape, location, scale, and gain. The first three are what control the parameters of the distribution and the last is a measure of the change in cars while traversing the road – to capture the effect of cars turning onto or off of the road on side streets.

Below, in Figure 2.23, an example estimate is performed. Two signals are provided (departure is pre-calculated within the agent and not shown here), trimmed, and then the optimization procedure is executed. The resulting distribution is overlaid on the downstream signal and estimated time delay gamma plotted above.

3.1.3 Updating Intra-Agent Behaviour Graphs

The second important piece in generating signal predictions is learning the intra-agent graph probabilities. These represent the likelihood of a signal to travel along this edge during the executed behaviours cycle.

3.1.3.1 Behaviour Probability Matrix (BPM) Figure 2.24 provides a visualization of this to help parse this idea. The matrix represents how these edge probabilities are used within the system; in the matrix,

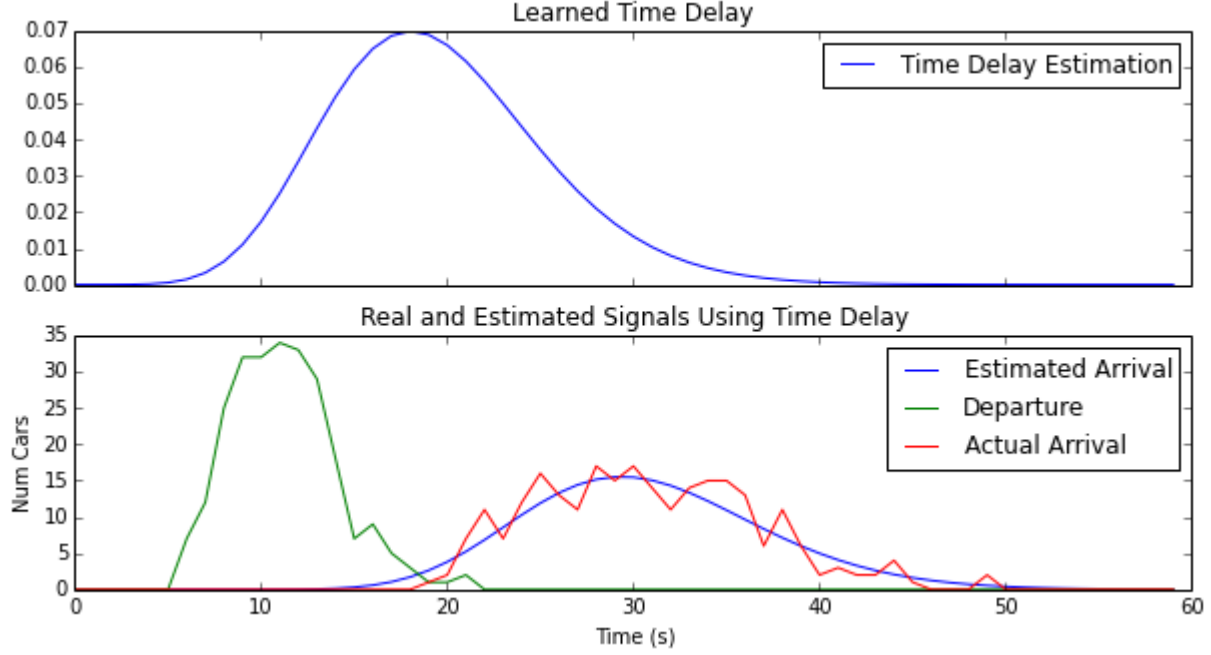


Figure 2.23: Example time delay estimation.

departing nodes are along the columns and entering along the rows – p_{02} would represent entering node 0 and exiting node 2. It is seen in the diagram that an event occurring at node 0 has a 60% likelihood of exiting node 2, for this behaviour. It is important to not that edges that are not traversable will always have a probability of 0. Essentially, this probability matrix is used to split an entering signal into components, to then be combined with other estimates in a departing signal. This estimated signal is then used in combination with the time delay estimation to predict expected traffic, which will be discussed further in the following section.

$$BPM = \begin{bmatrix} p_{00} & p_{01} & p_{02} & p_{03} \\ p_{10} & p_{11} & p_{12} & p_{13} \\ p_{20} & p_{21} & p_{22} & p_{23} \\ p_{30} & p_{31} & p_{32} & p_{33} \end{bmatrix}$$

3.1.3.2 Learning Probabilities

3.1.3.3 Updating BPMs After the optimization procedure has arrived at a solution, the existing BPM must be updated. It would be unwise to directly use the newly calculated BPM as temporary inhibitors to




Figure 2.24: Example intra-agent edge probabilities for node 0.

traffic flow, such as accidents, can greatly skew results. Therefore, a learning parameter is used to update existing probabilities, seen in Equation 2.1. This can be tuned, on an agent-by-agent basis, to alter how quickly these values are modified and adapt an average of the two values is a good starting point, to enable quick transitions, while remaining relatively stable. In 2.1, i represents the agents ID and BPM_{est}^i is the newly estimated behaviour probability matrix for that agent.

$$BPM_{new}^i = \alpha^i \times BPM_{old}^i + (1 - \alpha^i) \times BPM_{est}^i \quad (2.1)$$

t

Chapter 3

Design Testing and Validation

1.0 Front End

2.0 Back End

3.0 Deep End

Chapter 4

Recommended Design Modifications

Chapter 5

Timeline and Project Management

1.0 Conclusion